

Programming Assignment #1 (Due Feb 26 1 PM Central)

Problem: Implement a **fixed-depth decision tree algorithm**, that is, the input to the ID3 algorithm will include the training data and **maximum depth of the tree** to be learned. The code skeleton as well as data sets for this assignment can be found on e-Learning.

Data Sets: The data sets (in the folder `./data/`) are obtained from the UCI Repository and are collectively the MONK's Problem. These problems were the basis of a first international comparison of learning algorithms¹. The training and test files for the three problems are named `monks-X.train` and `monks-X.test`. There are six attributes/features (columns 2–7 in the raw files), and the class labels (column 1). There are 2 classes. Refer to the file `./data/monks.names` for more details.

- (**Learning Curves**, 40 points) For $\text{depth} = 1, \dots, 10$, learn decision trees and compute the average training and test errors on each of the three MONK's problems. **Make three plots, one for each of the MONK's problem sets**, plotting training and testing error curves together for each problem, with tree depth on the x -axis and error on the y -axis.
- (**Weak Learners**, 30 points) For `monks-1`, report the **learned decision tree** and the **confusion matrix on the test set** for `depth=1` and `depth=2`. A **confusion matrix** is a table that is used to describe the performance of a classifier on a data set. For binary classification problems, it will be:

		Classifier Prediction	
		Positive	Negative
Actual Value	Positive	True Positive	False Negative
	Negative	False Positive	True Negative

Figure 1: Confusion matrix for a binary classification problem.

- (**scikit-learn**, 15 points) For `monks-1`, use `scikit-learn`'s default decision tree algorithm² to learn a decision tree. Visualize the learned decision tree using `graphviz`³. Report the **visualized decision tree** and the **confusion matrix** on the test set. **Do not change the default parameters.**
- (**Other Data Sets**, 15 points) Repeat steps 2 and 3 with your “own” data set and report the confusion matrices. You can use other data sets in the UCI repository. If you encounter continuous features, consider a simple discretization strategy to pre-process them into binary features using the mean. For example, a continuous feature x can be discretized using its mean μ as

$$x_{\text{binary}} = \begin{cases} 0, & \text{if } x \leq \mu, \\ 1, & \text{if } x > \mu. \end{cases}$$

¹<https://archive.ics.uci.edu/ml/datasets/MONK's+Problems>

²<http://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>

³see <http://scikit-learn.org/stable/modules/tree.html#classification> for an example.